

# Applied Parallel Computing with Python – List of Tasks

**PyCon 2013**


# Goal

- Tackle CPU-bound tasks
- Accept the GIL
- Utilise many cores on many machines
- Maybe utilise many languages too

# Overview (pre-requisites)

- multiprocessing
- ParallelPython
- hotqueue, redis (and Redis system)
- Matplotlib (for visualisations)

# Mandelbrot as surface plot



# Serial single thread

- `$ python serial_python.py  
--plot3D --size 100`
- **2500 elements**
- `$ python serial_python.py`
- **250,000 elements**
- **11 seconds on 1 core**

# Amdahl's law

- Max speed-up is limited to the parallelisable portions and resources
- What serial constraints do we have?
- How many data elements?
- How much memory?
- What affects transmission speed?  
Gigabit? Switches? Traffic?

# Memory usage?

- `import sys`
- `sys.getsizeof(0+0j)` # 32 bytes
- `250,000 * 32 == ?` # lower-bound
- Pickling and sending will take time
- Assembling the result will take time

# Profile memory usage

- Github `fabianp memory_profiler`
- `$ python -m memory_profiler serial_python_temp.py #argparse`
- **Output (takes a while):**
- `61:q.append(complex...) # +25MB`
- `65:...=calculate_z(...) # +7MB`



# multiprocessing

- Using all our CPUs is cool, 4 are common, 32 will be common
- Global Interpreter Lock (isn't our enemy)
- Silo'd processes are easiest to parallelise
- Forks on local machine (1 machine only)
- <http://docs.python.org/library/multiprocessing>

# Making chunks of work

- Split the work into chunks
- Start splitting by number of CPUs
- Submit the jobs with `map_async`
- Get the results back, join the lists
- Profile and consider the results...

# multiprocessing Pool

- `2_mandelbrot_multiprocessing/`
- `multiproc.py`
- `p = multiprocessing.Pool()`
- `po = p.map_async(fn, args)`
- `result = po.get()` # for all po objects
- **join the result items to make full result**

# multiprocessing

- 1 process takes 12 secs
- 2 takes 6 secs (watch System Monitor)
- 4 takes about 5 – what's happening?
- What about 32?

# ParallelPython

- Same principle as multiprocessing but allows  $>1$  machine with  $>1$  CPU
- <http://www.parallelpython.com/>
- Seems to work poorly with lots of data (e.g. 8MB split into 4 lists...!)
- We can run it locally, run it locally via `ppserver.py` and run it remotely too
- Can we demo it to another machine?

# Running ParallelPython

- Run
- `$ python paralleipy.py #chunks`
- **Now to run server separately:**
- `$ ppserver.py -d -a # uses all CPUs`
- `$ python paralleipy_manymachines.py`

# ParallelPython + binaries

- We can ask it to use modules, other functions and our own compiled modules
- Works for Cython and ShedSkin
- Modules have to be in PYTHONPATH (or current directory for ppserver.py)

# “timeout: timed out”

- Beware the timeout problem, the default timeout isn't helpful:
  - `pptransport.py`
  - `TRANSPORT_SOCKET_TIMEOUT = 60*60*24 # from 30s`
- Remember to edit this on *all* copies of `pptransport.py`



# Redis queue

- Queue is persistent, architect. agnostic
- Server/client model, time shift ok
- 1\$ `python hotq.py # worker(s)`
- 2\$ `python hotq.py --server`
- What if many jobs get posted and you're consumers aren't running?
- Also->Amazon Simple Queue Service