

Applied Parallel Computing

Ian Ozsvald

ian@morconsulting.com

www.morconsulting.com

Minesh B. Amin

mamin@mbasciences.com


www.mbasciences.com

PyCON 2013

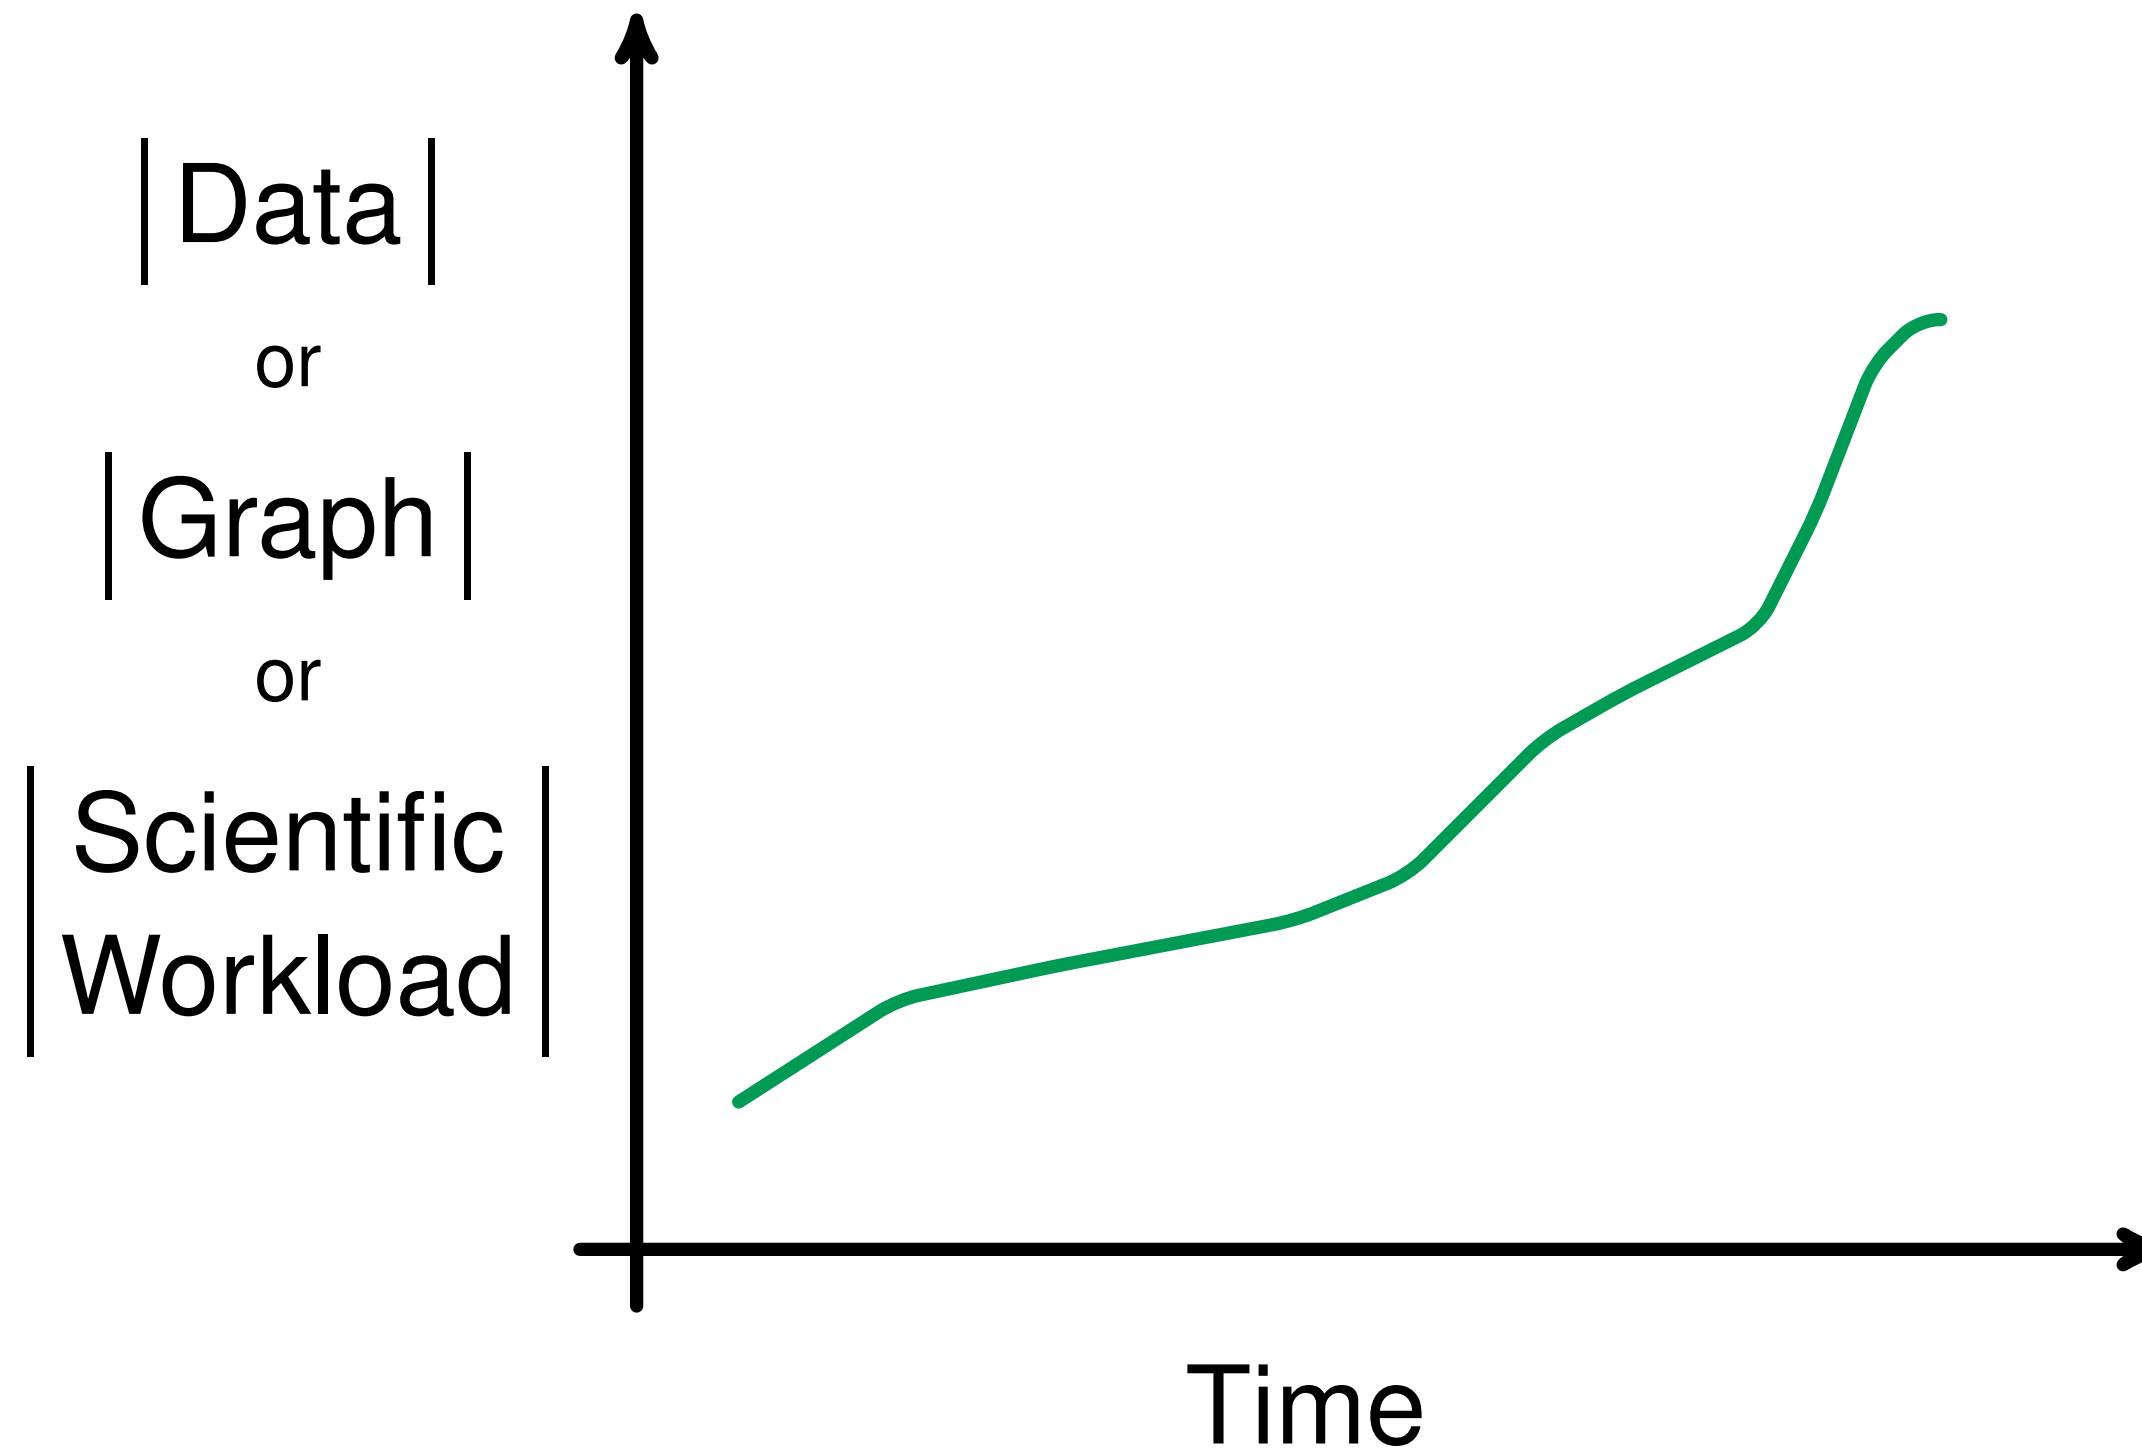
Santa Clara, CA

March 14, 2013

{ Robust Fault tolerant } Parallelism: Why?



{ Robust Fault tolerant } Parallelism: Why?



Parallelism: How (Take I)?

Serial Module

[illegible]

Parallelism: How (Take I)?

Serial Module

↓

Multiple

[illegible]

```
def taskEval(remoteArgs):
    # "cpg/ai4edparty/imageToEdges/b1a/run_12345-appar.sh"
    rankID = sgm.util.rankPolicy(0)

    return sgm.util.expressEvalShellPolicy(
        (cmd = "%(b1a)s" %
            %(id)s %
            %(dotConfig)s %
            %(dotType)s %
            %(rankID)s %
            %(dicFile)s % remoteArgs.id,"
            dotConfig = %(remoteArgs.dotConfig)s,
            dotType = %(remoteArgs.dotType)s,
            rankID = %(remoteArgs.rankID)s,
            seconds = %(remoteArgs.timeout)s);

    timeout = sgm.util.timeout.after
```

```
def taskval(remoteArgs):
    bin = "/opt/thirdparty/imageTools/bin/run_122NWrapper.sh";
    randM = spm.util.rand(policyM);
    return spm.util.compose.shell(policyM)
        (cmd = "%(bin)s" %
            %(id)s %
            %(dotConfig)s %
            %(dotType)s %
            %(randM)s" % dict(bin=bin, id=remoteArgs.id, dotConfig=remoteArgs.dotConfig, dotType=remoteArgs.dotType, randM=randM, timeout=remoteArgs.timeout));
    timeout = spm.util.timeout.after(seconds
```

Parallelism: How (Take I)?

Serial Module

Multiple Invocations

Parallel Module

[illegible]

```
def taskVal(remoteArgs):
    # = /opt/libdarty/imagetools/bin/run_IMAGEapp.sh'
    rank = sgm.util.prankPolicy()

    return sgm.util.sopposeEvalPolicy(
        (cmd = "%(task)s" % \
            "%(id)s" % \
            "%(destConfig)s" % \
            "%(destType)s" % \
            "%(rank)s" % sgmUtil,
         destConfig = remoteArgs.destConfig,
         destType = remoteArgs.destType,
         rank = rank,
         seconds = remoteArgs.timeout)),
        timeout = sgm.util.timeoutAfter
```

[illegible]

Parallelism: How (Take I)?

Serial Module

Multiple Invocations

Parallel Module

Multiple Invocations

```
def taskEval(remoteArgs):
    """ /opt/hadoop/bin/hadoop jar /usr/lib/hadoop/apper.sh """
    rank0 = sgm.util.rank0_policy()

    return sgm.util.cospecos(should_policy()
        (cmd = ["hadoop"],
            ["-idc"],
            ["-rankType"],
            ["-rank0"],
            dict(hs =
                dict(hs =
                    rank0Arg = remoteArgs.rank0,
                        doNotConfig = remoteArgs.doNotConfig,
                        doNotp = remoteArgs.doNotp,
                        rank0 = rank0,
                        rank0p = rank0p,
                        timeout = sgm.util.timeout_after(seconds = remoteArgs.timeout));
```

```
def task0Val(runcatArgs):
    ... = "/opt/7bitsec/imgs/7bitSec/bin/run_7bitSecWrapper.sh"
    runcat = sgm.util.prank.policy(0)


    return sgm.util.sopagate.shell.policy(0)
        (cmd) = "%lsblk" ~ "\n"
            ~ "%lsblk" ~ "\n"
            ~ "%dotConfig" ~ "\n"
            ~ "%dotTgr" ~ "\n"
            ~ "%rand" ~ "\n"

        if (0) == runcatArgs.i.d,
            ~ runcatArgs.dotConfig,
            ~ runcatArgs.dotTgr,
            ~ rand,
            ~ runcatArgs.timeout))

timeout = sgm.util.timeout.after
        (seconds = runcatArgs.timeout))
```

[illegible][illegible]

Robust Fault tolerant Parallelism: How (Take I)?



Robust Fault tolerant Parallelism: How (Take I)?

Serial
Module

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

Multiple
Invocations

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

Parallel
Module

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

Multiple
Invocations

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

Parallel
Workflow

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

Multiple
Invocations

```
def taskEval(remoteArgs):
    rankD = sgm.util.prank.policy()
    return sgm.util.coprocess.shell.policy()
    (cmd = "%(bin)s" % \
        "%(dotConfig)s" % \
        "%(rankD)s" % dict(bin = remoteArgs.id,
                             dotConfig = remoteArgs.dotConfig,
                             dotType = remoteArgs.dotType,
                             rankD = rankD,
                             seconds = remoteArgs.timeout))
    timeout = sgm.util.timeout.after(seconds = remoteArgs.timeout))
```

{ Robust Fault tolerant } Parallelism: How (Take II)?

A Language determines the concepts we can think of
- Benjamin Worf

$\left\{ \begin{array}{l} \text{Robust} \\ \text{Fault tolerant} \end{array} \right\}$ Parallelism: How (Take II)?

A $\left\{ \begin{array}{l} \text{Language} \end{array} \right\}$ determines the concepts we can think of

$\left\{ \begin{array}{l} \text{Robust} \\ \text{Fault tolerant} \end{array} \right\}$ Parallelism: How (Take II)?

A $\left\{ \begin{array}{l} \text{Language} \\ \text{Runtime Env} \end{array} \right\}$ determines the concepts we can think of

$\left\{ \begin{array}{l} \text{Robust} \\ \text{Fault tolerant} \end{array} \right\}$ Parallelism: How (Take II)?

A $\left\{ \begin{array}{l} \text{Language} \\ \text{Runtime Env} \\ \text{Framework} \end{array} \right\}$ determines the concepts we can think of

$\left\{ \begin{array}{l} \text{Robust} \\ \text{Fault tolerant} \end{array} \right\}$ Parallelism: How (Take II)?

A $\left\{ \begin{array}{l} \text{Language} \\ \text{Runtime Env} \\ \text{Framework} \\ \text{Library} \end{array} \right\}$ determines the concepts we can think of

Preamble: "The Big Picture"

Question: Is exploiting parallelism $\left\{ \begin{array}{l} \text{easy} \\ \text{hard} \end{array} \right\}$?

Preamble: "The Big Picture"

What makes

Question: ~~Is~~ exploiting parallelism $\left\{ \begin{array}{l} \text{easy} \\ \text{hard} \end{array} \right\}$?



Copyright 1994, The UNIX-HATERS Handbook

Preamble: "The Big Picture"

What makes

Question: ~~Is~~ exploiting parallelism $\left\{ \begin{array}{l} \text{easy} \\ \text{hard} \end{array} \right\}$?



Copyright 1994, The UNIX-HATERS Handbook

Supposition: The gap between developer's intent and API of PET
(parallel enabling technologies) ...


Preamble: "Parallel Enabling Technologies"

Means to the end

- Bottom-up

OpenMPI OpenMP
CUDA OpenGL

- Maximum flexibility
- Maximum headaches
- Must implement fault tolerance




Preamble: "Parallel Enabling Technologies"

Means to the end

- Bottom-up

OpenMPI OpenMP
CUDA OpenGL


- Maximum flexibility
- Maximum headaches
- Must implement fault tolerance



- Top-down

Hadoop Goldenorb
GraphLab

- Limited flexibility
- Fewer headaches
- Fault tolerance is inherited




Preamble: "Parallel Enabling Technologies"

Means to the end

- Bottom-up

OpenMPI OpenMP
CUDA OpenGL


- Maximum flexibility
- Maximum headaches
- Must implement fault tolerance



- Top-down

Hadoop Goldenorb
GraphLab


- Limited flexibility
- Fewer headaches
- Fault tolerance is inherited



- Self-contained environment

SPM.Python

- Maximum flexibility
- Fewest headaches
- Fault tolerance is inherited



Preamble: "Exploiting Parallelism"

Parallelism: The management of a collection of serial tasks

Management: The policies by which:

- tasks are scheduled,
- premature terminations are handled,
- preemptive support is provided,
- communication primitives are enabled/disabled, and
- the manner in which resources are obtained and released

Serial Tasks: Are classified in terms of either:

- *Coarse Grain* ... where tasks may not communicate prior to conclusion, or
- *Fine Grain* ... where tasks may communicate prior to conclusion.

Preamble: "Exploiting Parallelism"

Parallelism: The management of a collection of serial tasks

Management: The policies by which:

- tasks are scheduled,
- premature terminations are handled,
- preemptive support is provided.

Management policies codify **how** serial tasks are to be managed ... independent of **what** they may be


Serial Tasks: Are classified in terms of either:

- **Coarse Grain** ... where tasks may not communicate prior to conclusion, or
- **Fine Grain** ... where tasks may communicate prior to conclusion.

management:

... a more challenging facet of [parallel] software engineering ...

- The Future of Computing Performance: Game Over or Next Level?
National Academy of Sciences, 2011



Module (Serial)

Module (Parallel)

Parallel Workflow

```
def taskEval(remoteArgs):
    return spm.util.coprocess.shell.policyB\
        (cmd = "%(bin)s" \
             "%(v)s" \
             "-c %(c)s" \
             "-d %(d)s" \
             "-s %(s)s" \
             "-a %(a)s" \
             "-o %(o)s" % dict(bin = remoteArgs.bin,
                               v = { True : '-v',
                                     False : '',
                                     None : '' },
                               c = remoteArgs.c,
                               d = remoteArgs.d,
                               s = remoteArgs.s,
                               a = remoteArgs.a,
                               o = remoteArgs.o,
                               ),
          timeout = spm.util.timeout.after(seconds = remoteArgs.timeout));
```

Module (Serial)

Module (Parallel)

Parallel Workflow

```
def taskEval(remoteArgs):  
  
    def taskEval(remoteArgs):  
        bin = "/opt/thirdparty/imageToEdge/bin/run_I2EAWrapper.sh";  
        rankD = spm.util.prank.policyD();  
  
        return spm.util.coprocess.shell.policyB\  
            (cmd = "%(bin)s" \\  
              "%(id)s" \\  
              "%(dotConfig)s" \\  
              "%(dotTgz)s" \\  
              "%(rankD)s" % dict(bin = bin,  
                                id = remoteArgs.id,  
                                dotConfig = remoteArgs.dotConfig,  
                                dotTgz = remoteArgs.dotTgz,  
                                rankD = rankD,  
                                ),  
              timeout = spm.util.timeout.after(seconds = remoteArgs.timeout));  
  
    timeout = spm.util.timeout.after(seconds = remoteArgs.timeout);
```

Module (Serial)

Module (Parallel)

Parallel Workflow

```
def main(pool, env):
    submitTask(env = env) \
    .execTask (pool = pool,
               timeoutWaitForSpokes = 2, # Secs
               timeoutExecution = 300, # Secs
               );

    if (terminateEarly):
        raise Exception("skysim failed");

@grainCoarseSingleton.pclosure
def submitTask():
    return stdlib.cache(# Core options
                        bin = '...',
                        v = True,
                        c = "/nfs/expt100000/dotCConfig.json",
                        d = "/nfs/expt100000/dotD",
                        s = "/nfs/expt100000/dotSConfig.json",
                        a = "/nfs/expt100000/dotAConfig.json",
                        o = "/nfs/expt100000/output",
                        # Meta options
                        timeout = 300, # Secs
                        label = "phaseA (exec)",
                        env = env);
```

Module (Serial)

Module (Parallel)

Parallel Workflow

```
def main(pool, env, batchFile):
    submitTask(env, batchFile) \
    .execTask(pool,
               timeoutWaitForSpokes = 2, # Secs
               timeoutExecution      = 300, # Secs
               );
```

```
@grainCoarseList.pclosure
def submitTask(env, batchFile):
    rval = [];
    for cmd in filter(len,                # Skip any empty line ...
                      map((lambda x:
                           x.strip()),   # Skip any prefix/suffix spaces ...
                          batchFile.split('\n'))):
        rval += [ stdlib.cache(cmd = cmd,
                                timeout = 2, # Secs
                                label    = "imageToEdge (exec - %(ct)d)" \
                                           dict(ct = len(rval),)),
                  ],
    return rval;
```

Module (Serial)

Module (Parallel)

Parallel Workflow

```
def main():
    import __hidden__.pool as pool;
    import __hidden__.env as env;
    import util.phaseA.par as phaseA;
    import util.phaseB.par as phaseB;
    import util.phaseC.par as phaseC;
    import util.phaseD.par as phaseC;

    try:
        pool = pool.interAll();
        env = env.main();

        phaseA.main(pool = pool, env = env);
        phaseB.main(pool = pool, env = env);
        batchFile = phaseC.main(pool = pool, env = env);
        phaseD.main(pool = pool, env = env, batchFile = batchFile);
    except Exception e:
        ...

    return;
```

Module (Serial)

Module (Parallel)

Parallel Workflow

```
def main():
```

```
    def main():
```

```
        import __hidden__.pool as pool;
```

```
        import __hidden__.env as env;
```

```
        import util.mc .par as mc;
```

```
    try:
```

```
        pool = pool.interAll();
```

```
        env = env.main();
```

```
        mc.main(pool = pool, env = env);
```

```
    except Exception e:
```

```
        ...
```

```
    return;
```

```
env = env);
```

```
env = env);
```

```
env = env);
```

```
env = env, batchFile = batchFile);
```

```
except Exception e:
```

```
...
```

```
return;
```

Module (Parallel): Intra-node

Device-specific Component
(in Emerald, C, C++ and Fortran)

```
def::api pow2::void <Target = Cuda> \
    (var a::matrixA&):
    # wrapper around implementation in C++

def::api pow2::void <Target = X98Cores> \
    (var a::matrixA&):
    # wrapper around implementation in C
    ...

def::api pow2::void <Target = Serial> \
    (var a::matrixA&):
    ...
```

Module (Parallel): Intra-node

Heterogeneous Component (in Emerald)

```
def::api main::void (dim1::int, dim2::int):  
    var a::matrixA[dim1,dim2] = rand;  
    var b::matrixB[dim1,dim2] = rand;  
  
    try::concurrent:  
        from ( demo :: explicit ) import pow2;  
  
        pow2(a);  
        b *= 2.0;  
    except:  
        raise;  
  
    assert(a::(Cuda == Serial == X86Cores));  
    assert(b::(Cuda == Serial));  
  
    from global import result;  
  
    result = a::Cuda;  
  
    return;
```

Device-specific Component (in Emerald, C, C++ and Fortran)

```
::void <Target = Cuda> \  
    (var a::matrixA&):  
    round implementation in C++  
  
::void <Target = X98Cores> \  
    (var a::matrixA&):  
    round implementation in C  
  
::void <Target = Serial> \  
    (var a::matrixA&):
```


Module (Parallel): Intra-node

Heterogeneous Data Structure (in Emerald)

```
def::struct myMatrix (nDim1 :: int,  
                      nDim2 :: int):  
  
  - @ -::Array  
    Domain = (nDim1, nDim2);  
    Format = Row;  
    Target = (::CUDA, ::Serial, ::X86Cores);  
  
  - @ -  
  {  
    float _;  
  };
```

Global Component (in Emerald)

```
def (dim1::int, dim2::int):  
  a[dim1,dim2] = rand;  
  b[dim1,dim2] = rand;  
  
  from (demo :: explicit) import pow2;  
  
  pow2(a);  
  b *= 2.0;  
  except:  
    raise;  
  
  assert(a:: (Cuda == Serial == X86Cores));  
  assert(b:: (Cuda == Serial));  
  
  from global import result;  
  
  result = a::Cuda;  
  
  return;
```

Device-specific Component (in Emerald, C, C++ and Fortran)

```
::void <Target = Cuda> \  
  (var a::matrixA&):  
    round implementation in C++  
  
::void <Target = X98Cores> \  
  (var a::matrixA&):  
    round implementation in C  
  
::void <Target = Serial> \  
  (var a::matrixA&):
```

Suppositions

Most embarrassingly parallel solutions
perform a lot of redundant work ...

Suppositions

Most embarrassingly parallel solutions
perform a lot of redundant work ...



- Only fix ... share knowledge

Suppositions

Many problems in HPC and Analytics are memory bounded ...

Suppositions

Many problems in HPC and Analytics are memory bounded ...



- Cannot depend on virtualization
- Must throw everything at the problem

Suppositions: Constraints When Exploiting Parallelism

Prototypes
vs
Runtime Env

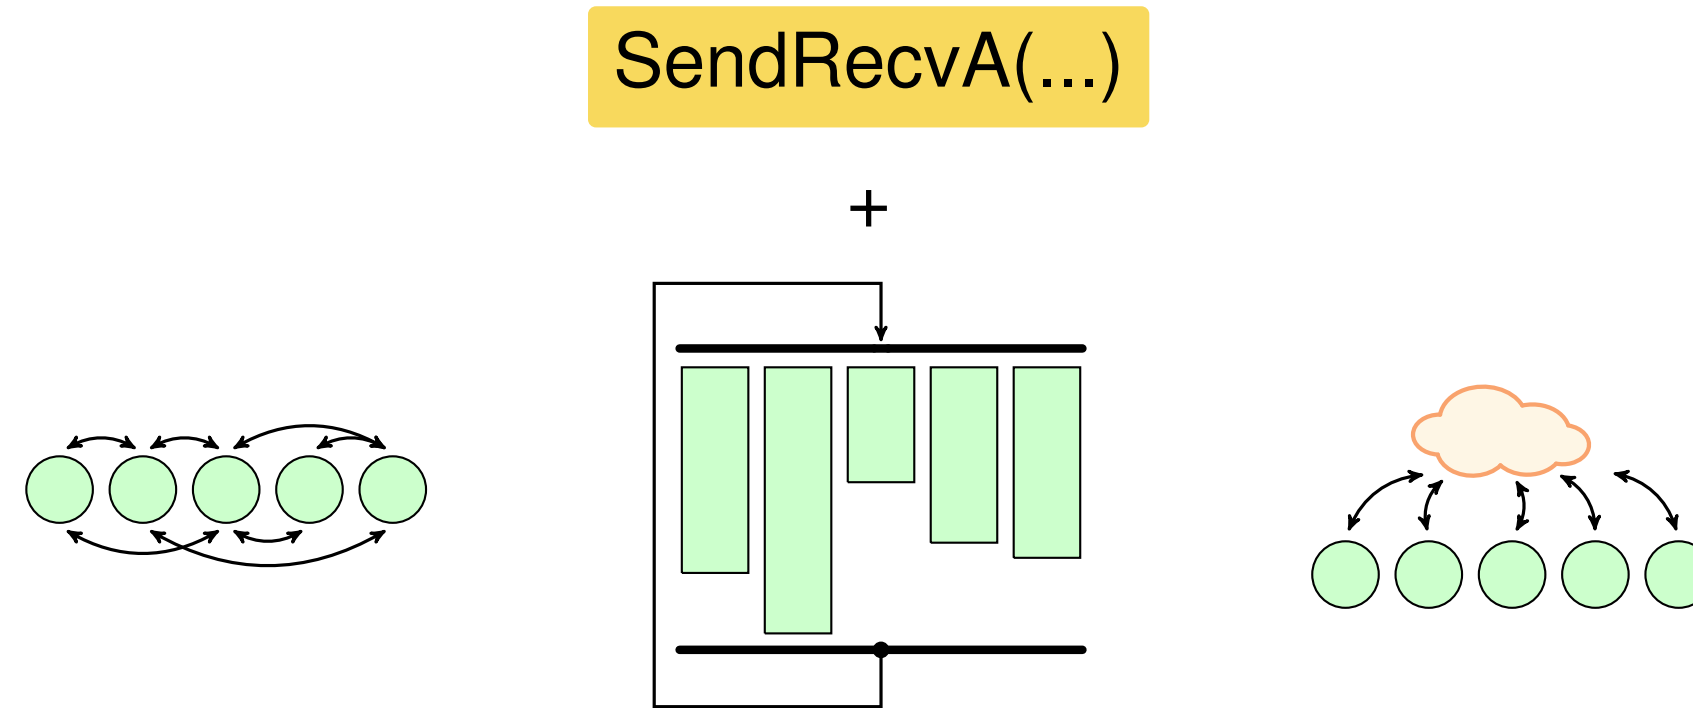
Suppositions: Constraints When Exploiting Parallelism

Prototypes → Construct-by-correction
vs
Runtime Env


Suppositions: Constraints When Exploiting Parallelism

Prototypes	→	Construct-by-correction
vs		
Runtime Env	→	Correct-by-construction


Suppositions: Enabling / Disabling Communication Primitives




Suppositions: Enabling / Disabling Communication Primitives




Suppositions: Enabling / Disabling Communication Primitives




Suppositions: Enabling / Disabling Communication Primitives



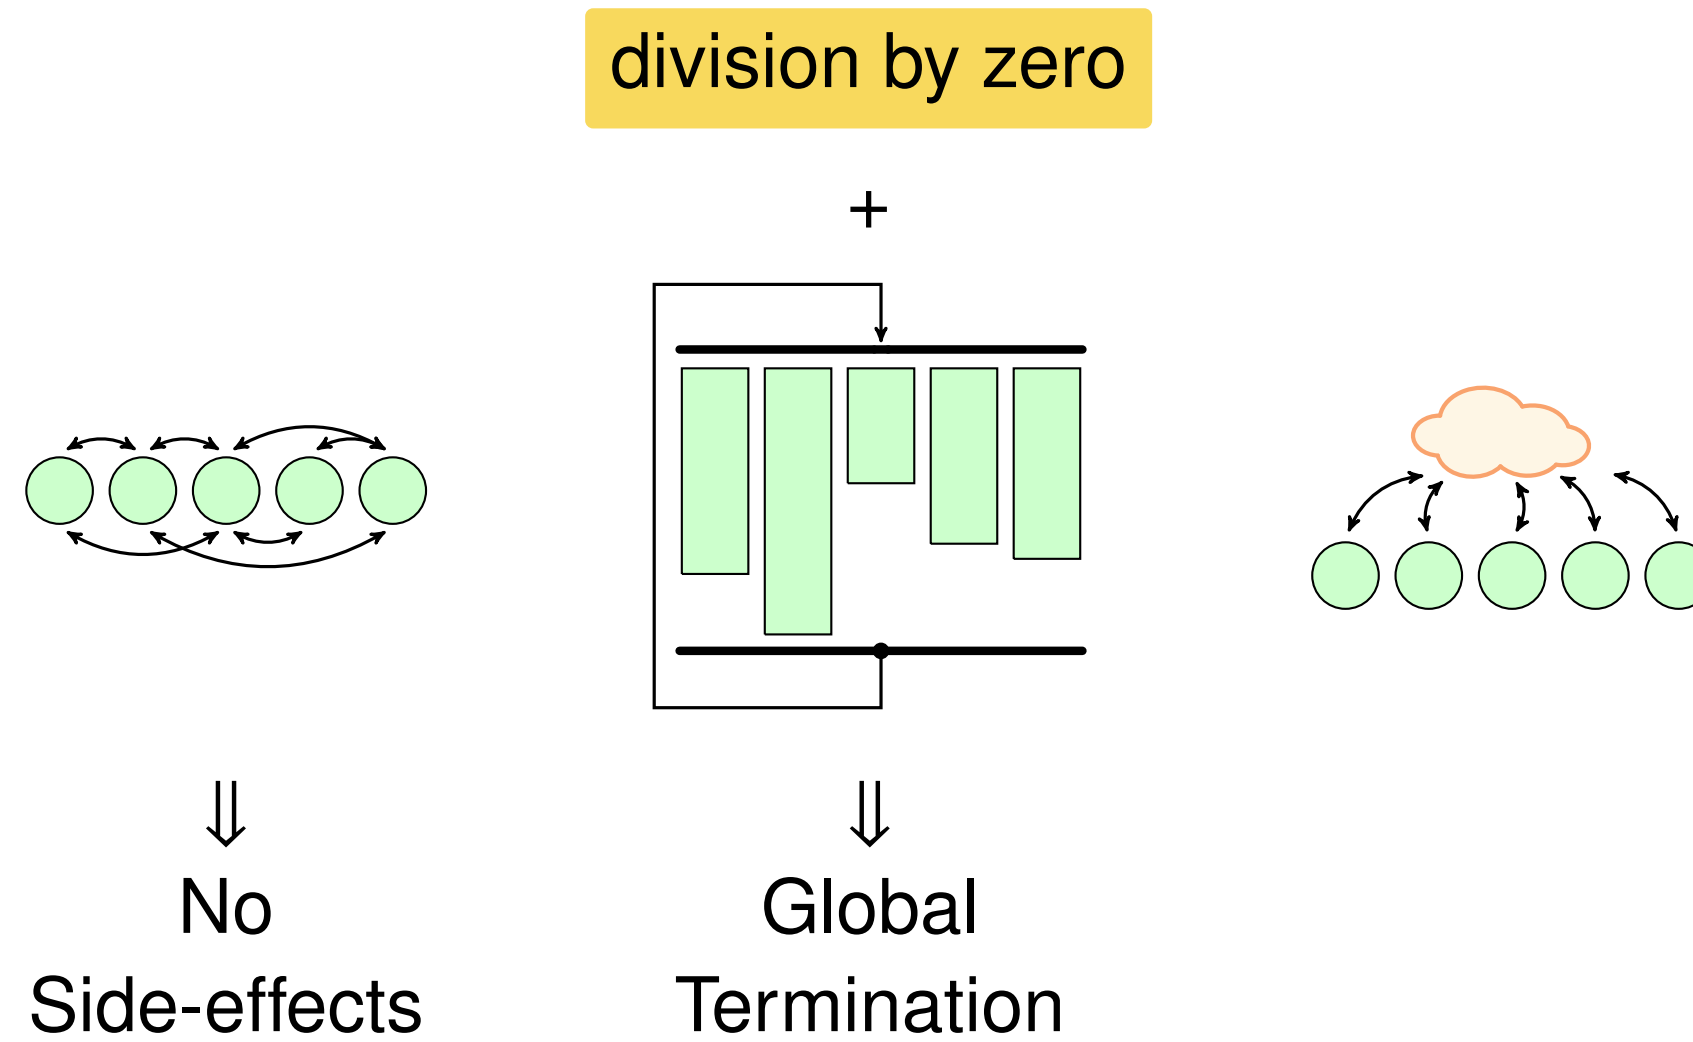
Suppositions: Parallel Semantics



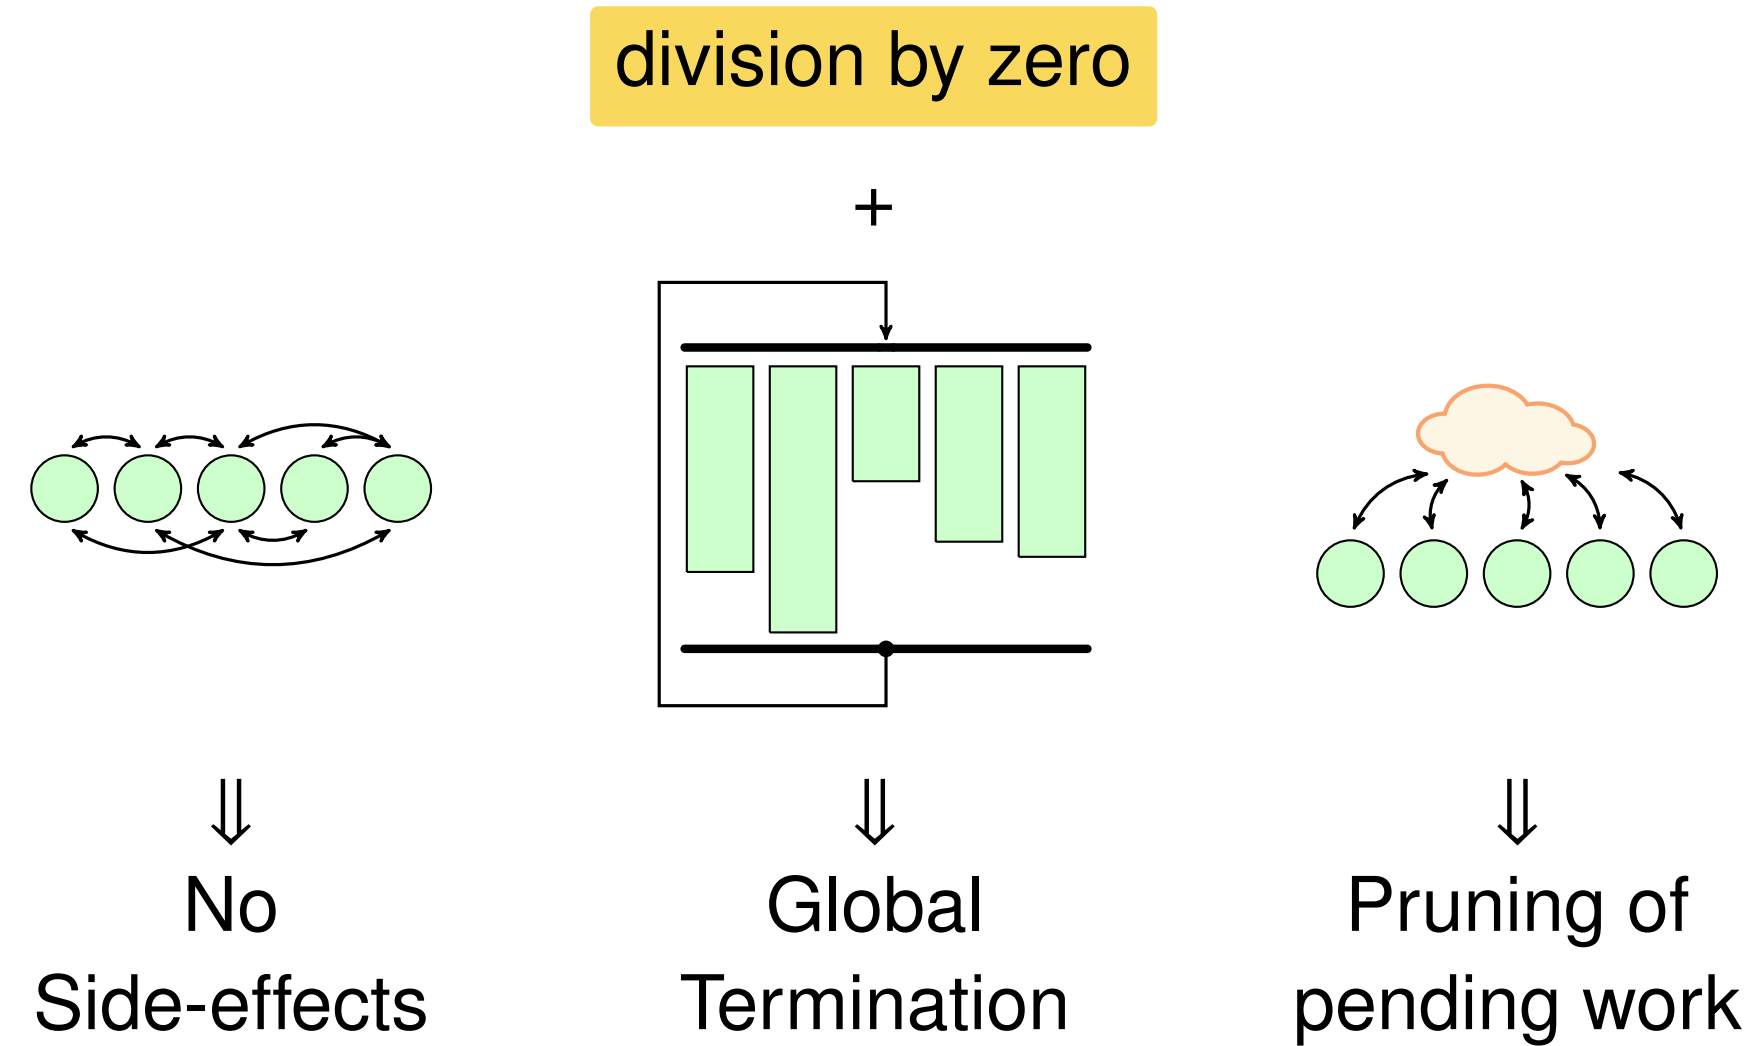
Suppositions: Parallel Semantics



Suppositions: Parallel Semantics



Suppositions: Parallel Semantics



Conclusion: Rest of the tutorial

For each form of parallelism to be reviewed:

- What is the management policy?
- Describe a compatible communication primitive
- Describe a toxic communication primitive